

Lecture 1: Computer Architecture Introduction

Today's topics:

เรียน CPU : MIPS

- Why computer organization is important
- Logistics
- Modern trends

Why Computer Organization




Image credits: uber, extremetech, anandtech

Why Computer Organization



Why Computer Organization

Embarrassing if you are a BS in CS/CE and can't make sense of the following terms: DRAM, pipelining, cache hierarchies, I/O, virtual memory, ...

Embarrassing if you are a BS in CS/CE and can't decide which processor to buy: 4.4 GHz Intel Core i9 or 4.7 GHz AMD Ryzen 9 (reason about performance/power)

Obvious first step for chip designers, compiler/OS writers

Will knowledge of the hardware help you write better and more secure programs?

Must a Programmer Care About Hardware?

Must know how to reason about program performance and energy and security

Memory management: if we understand how/where data is placed, we can help ensure that relevant data is nearby

มีทั้งแบบ HW, SW

Thread management: if we understand how threads interact, we can write smarter multi-threaded programs

□ Why do we care about multi-threaded programs?

Example

เวลาเก่า/เวลาใหม่

200x speedup for matrix vector multiplication

Data level parallelism: 3.8x

คูณกันทั้งหมดได้ประมาณ 200

Loop unrolling and out-of-order execution: 2.3x

HW

HW Cache blocking: 2.5x

Thread level parallelism: 14x

GPU

Further, can use accelerators to get an additional 100x.

Key Topics

จำนวนทรานซิสเตอร์บน Chip จะเพิ่มขึ้นสองเท่าทุกๆ 18-24 เดือน

Moore's Law, power wall

Use of abstractions

MIPS Assembly language

~~Computer arithmetic~~

Pipelining

Using predictions

Memory hierarchies

Accelerators

Reliability and Security

Logistics

ไปดูคลิปของเค้ามา สำคัญมาก!

See class web-page

<http://www.cs.utah.edu/~rajeev/cs3810>

TAs and office hours: TBA

Most communication on Canvas; email me directly to set up office hours, or meet me right after class

Textbook: Computer Organization – HW/SW Interface, Patterson and Hennessy, 5th edition

Course Organization

30% midterm, 30% final, 30% Lab+assign, 10%Atten


Co-operation policy: you may discuss – you may not see someone else's written matter when writing your solution

Exams are open-book and open-notes, Open-Internet

Print slides just before class

Screencast YouTube videos

Microprocessor Performance



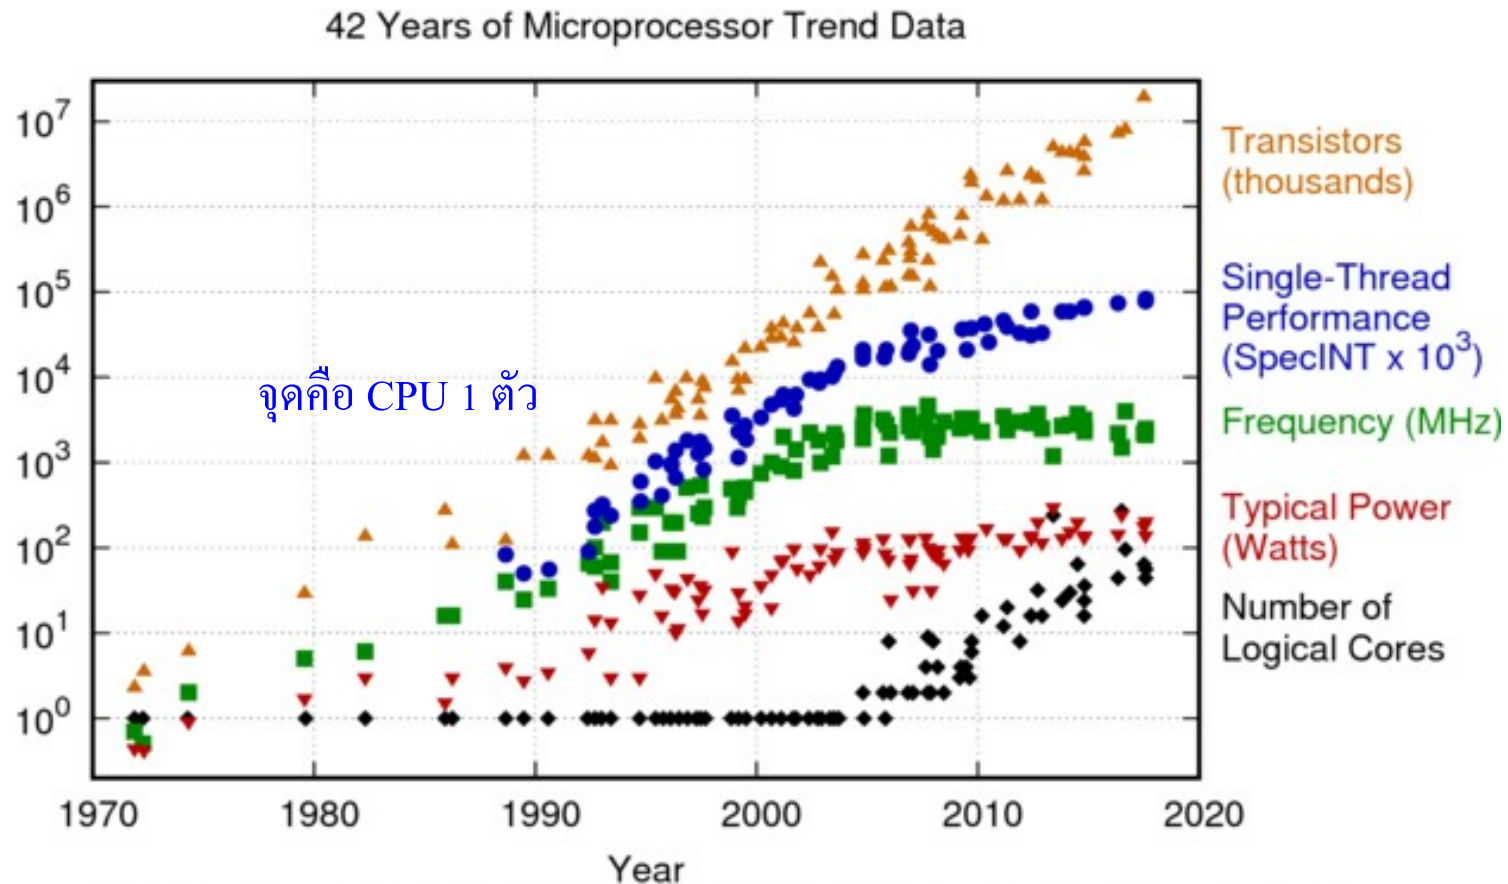
ปี 2003 เราติดปัญหา Powerwall และวิธีต่างๆถูกค้นพบไปหมดแล้ว

50% improvement every year!!

What contributes to this improvement?

Source: H&P Textbook


Microprocessor Performance



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

Source: karlrupp.net

Microprocessor Performance



Source: karlrupp.net

Power Consumption Trends

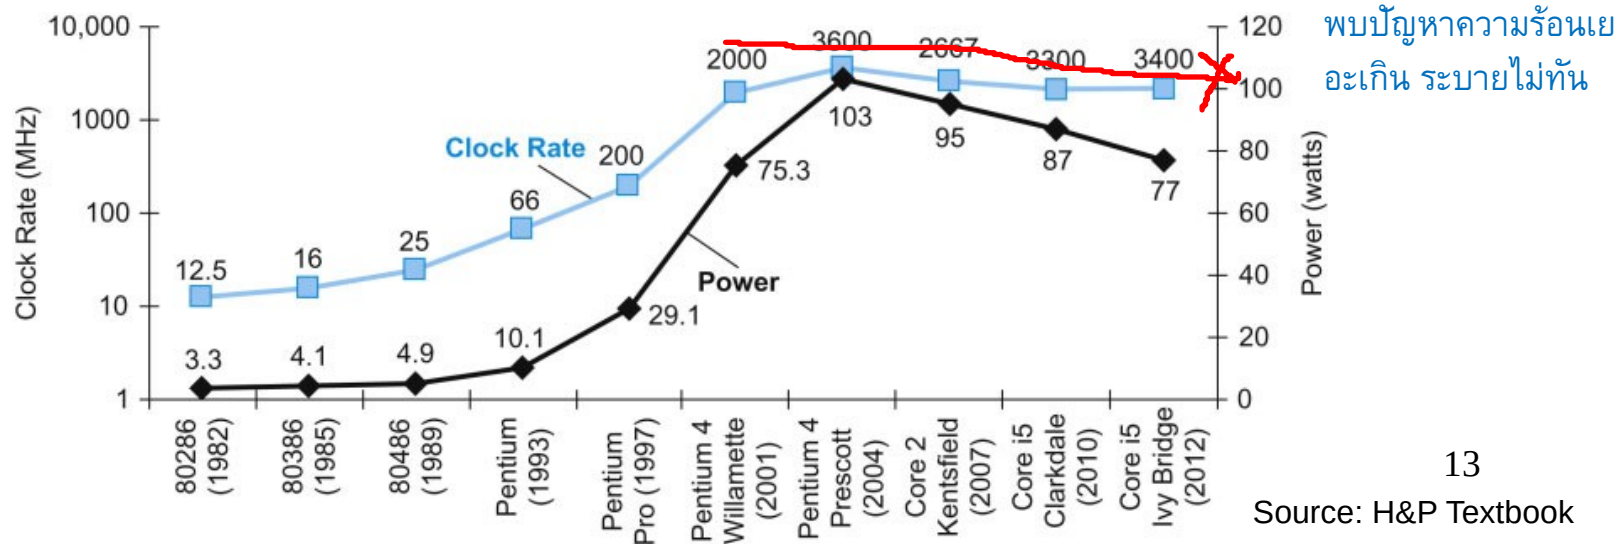
สมการนี้สำคัญมาก

แปรผันตรง
Dyn power \propto activity x capacitance x voltage² x frequency
Dynamic Power จำนวนทรานซิสเตอร์

Voltage and frequency are somewhat constant now,
while capacitance per transistor is decreasing and number
of transistors (activity) is increasing

เพราะว่า Transistor
มีขนาดเล็กลง

Leakage power is also rising (function of #trans and voltage)



Summary

สรุป สำคัญทุกตัว

Increasing frequency led to power wall in early 2000s

Frequency has stagnated since then

End of voltage (Dennard) scaling in early 2010s

Has led to dark silicon and dim silicon (occasional turbo)

บางคอร์ไม่ทำงานทำงานเมื่อต้องการ

Important Trends

Running out of ideas to improve single thread performance

Power wall makes it harder to add complex features

Power wall makes it harder to increase frequency

Additional performance provided by: more cores, occasional spikes in frequency, accelerators

GPU

Important Trends

Historical contributions to performance:

1. Better processes (faster devices) ~20%
2. Better circuits/pipelines ~15%
3. Better organization/architecture ~15% เราเลยเรียนวิชานี้

In the future, bullet-2 will help little and bullet-1 will eventually disappear! เพราะอนาคต Transister จะไม่เร็วขึ้นแล้ว

	Pentium	P-Pro	P-II	P-III	P-4	Itanium	Montecito
Year	1993	95	97	99	2000	2002	2005
Transistors	3.1M	5.5M	7.5M	9.5M	42M	300M	1720M
Clock Speed	60M	200M	300M	500M	1500M	800M	1800M

Moore's Law in action

At this point, adding transistors to a core yields little benefit

What Does This Mean to a Programmer?

Today, one can expect only a 20% annual improvement; the improvement is even lower if the program is not multi-threaded

- A program needs many threads
- The threads need efficient synchronization and communication
- Data placement in the memory hierarchy is important
- Accelerators should be used when possible

Challenges for Hardware Designers

Find efficient ways to

- improve single-thread performance and energy
- improve data sharing
- boost programmer productivity
- manage the memory system
- build accelerators for important kernels
- provide security

The HW/SW Interface

Application software

Systems software
(OS, compiler)

Hardware

$a[i] = b[i] + c;$

↓ Compiler

```
lw    $15, 0($2)
add   $16, $15, $14
add   $17, $15, $13
lw    $18, 0($12)
lw    $19, 0($17)
add   $20, $18, $19
sw    $20, 0($16)
```

↓ Assembler

```
000000101100000
110100000100010
...
```

Computer Components


Input/output devices

Secondary storage: non-volatile, slower, cheaper (HDD/SSD)

Primary storage: volatile, faster, costlier (RAM)

CPU/processor (datapath and control)

Wafers and Dies



Source: H&P Textbook

Manufacturing Process

Silicon wafers undergo many processing steps so that different parts of the wafer behave as insulators, conductors, and transistors (switches)

Multiple metal layers on the silicon enable connections between transistors

The wafer is chopped into many dies – the size of the die determines yield and cost

Processor Technology Trends

Shrinking of transistor sizes: 250nm (1997) □
130nm (2002) □ 70nm (2008) □ 35nm (2014) □
2019, start of transition from 14nm to 10nm

Transistor density increases by 35% per year and die size increases by 10-20% per year... functionality improvements!

Transistor speed improves linearly with size (complex equation involving voltages, resistances, capacitances)

Wire delays do not scale down at the same rate as transistor delays

Memory and I/O Technology Trends

DRAM density increases by 40-60% per year, latency has reduced by 33% in 10 years (the memory wall!), bandwidth improves twice as fast as latency decreases

Disk density improves by 100% every year, latency improvement similar to DRAM

Networks: primary focus on bandwidth; 10Mb << 100Mb in 10 years; 100Mb << 1Gb in 5 years